

CSE 374 Programming concepts and tools

Winter 2024

Instructor: Alex McKinney



HW9



HW9: Concurrency

I've decided to make HW9 an **extra credit** assignment.

- Released on Sunday, March 3
- Due Sunday, March 10
 - Late days will not be accepted on this assignment.
 - No extensions.



This is your chance to recuperate some points you may have lost.

This assignment is challenging.

- It covers concurrency, which we will go over in class on Friday.
- It will require a substantial amount of **self-learning and research**.
- No autograder.

HW9: Concurrency (cont.)

The purpose of this assignment is to show you how far you've come!

- We've already learned three different programming languages thus far (bash, C/C++).
- You now have the skills required to learn any programming language out there.
- "Learn how to acquire additional information and skills independently" ([syllabus](#))



You will need to write code in a foreign code base ([Go](#)).

- TAs are not expected to have finished this assignment.
- **Assistance will be scarce.**

A Tour of Go

<https://go.dev/tour>

Go is modern language widely used in the software engineering industry (just like bash and C/C++).

- Open-sourced by Google
- Designed by C++ experts
- Syntax is similar to C, but has garbage collection.



```
hello.go Imports off Syntax off
1 package main
2
3 import "fmt"
4
5 func main() {
6     fmt.Println("Hello, 世界")
7 }
8
```

Reset Format Run

Questions?



Review

Review: `std::unique_ptr`

A `unique_ptr` **takes ownership** of a pointer

- Part of C++'s standard library (C++11)
- Its destructor invokes `delete` on the owned pointer
 - Invoked when `unique_ptr` object is `delete`'d or falls out of scope via the `unique_ptr` destructor
- Guarantees uniqueness by disabling copy and assignment.

Review: `std::shared_ptr`

`shared_ptr` is similar to `unique_ptr` but we allow shared objects to have multiple owners

- The copy/assign operators are not disabled and they *increment* **reference counts** as needed
- When a `shared_ptr` is destroyed, the reference count is decremented
 - When the reference count hits 0, we **delete** the pointed-to object!
- Allows us to create complex linked structures (double-linked lists, graphs, etc.) at the cost of maintaining reference counts

Review: `std::weak_ptr`

`weak_ptr` is similar to a `shared_ptr` but doesn't affect the reference count

- Can *only* “point to” an object that is managed by a `shared_ptr`
- Not *really* a pointer – can't actually dereference unless you “get” its associated `shared_ptr`
- Because it doesn't influence the reference count, `weak_ptr`s can become “*dangling*”
 - Object referenced may have been `delete`'d
 - But you can check to see if the object still exists
- Can be used to break our cycle problem!

Questions?



C++ Inheritance

Overview

C++ Inheritance

- Review of basic idea (pretty much the same as in Java)
- What's different in C++ (compared to Java)
 - *Static vs dynamic dispatch - virtual functions and vtables (i.e., dynamic dispatch) are optional*
 - *Pure virtual functions, abstract classes, why no Java “interfaces”*
 - *Assignment slicing, using class hierarchies with STL*
- HW8 Walkthrough

Motivation

and C++ Syntax

Stock Portfolio Example

A portfolio represents a person's financial investments

- Each *asset* has a cost (*i.e.* how much was paid for it) and a market value (*i.e.* how much it is worth)
 - The difference between the cost and market value is the *profit* (or loss)
- Different assets compute market value in different ways
 - A **stock** that you own has a ticker symbol (e.g. “GOOG”), a number of shares, share price paid, and current share price
 - A **dividend stock** is a stock that also has dividend payments
 - **Cash** is an asset that never incurs a profit or loss

Design Without Inheritance

One class per asset type:

Stock
symbol_ total_shares_ total_cost_ current_price_
GetMarketValue() GetProfit() GetCost()

DividendStock
symbol_ total_shares_ total_cost_ current_price_ dividends_
GetMarketValue() GetProfit() GetCost()

Cash
amount_
GetMarketValue()

- Redundant!
- Cannot treat multiple investments together
 - e.g. can't have an array or vector of different assets

Inheritance

A parent-child “is-a” relationship between classes

- A child (**derived class**) extends a parent (**base class**)

Terminology:

Java	C++
Superclass	Base Class
Subclass	Derived Class


- Java: Subclass inherits from super class. (Superclass is “higher” in the hierarchy)
- C++: Derived class inherits from base class. (Base class is “higher” in the hierarchy)
 - Think of derived class as a derivative of the base class (e.g. car class is a derivative of vehicle class)
 - Mean the same things. You’ll hear both.

Inheritance

Benefits:

- Code reuse
 - Children can automatically inherit code from parents
- Polymorphism
 - Ability to redefine existing behavior but preserve the interface
 - Children can override the behavior of the parent
 - Others can make calls on objects without knowing which part of the inheritance tree it is in
- Extensibility
 - Children can add behavior

Design With Inheritance



Access Modifiers

- `public:` visible to all other classes
- `protected:` visible to current class and its *derived* classes
- `private:` visible only to the current class

Use `protected` for class members only when

- Class is designed to be extended by subclasses
- Derived classes must have access but clients should not be allowed

`protected` isn't truly protected as an adversary client can just extend a protected class and get access to the "protected" information. Hence its use is rather limited in the real world.

Class derivation List

Comma-separated list of classes to inherit from:

```
#include "BaseClass.h"
class Name : public BaseClass {
    ...
};
```

- Focus on **single inheritance**, but *multiple inheritance* possible
 - : public Base1, public Base2 {

Almost always you will want **public inheritance**

- Acts like `extends` does in Java
- Any member that is non-private in the base class is the same in the derived class; both *interface and implementation inheritance*
 - Except that constructors, destructors, copy constructor, and assignment operator are *never* inherited

Back to Stocks

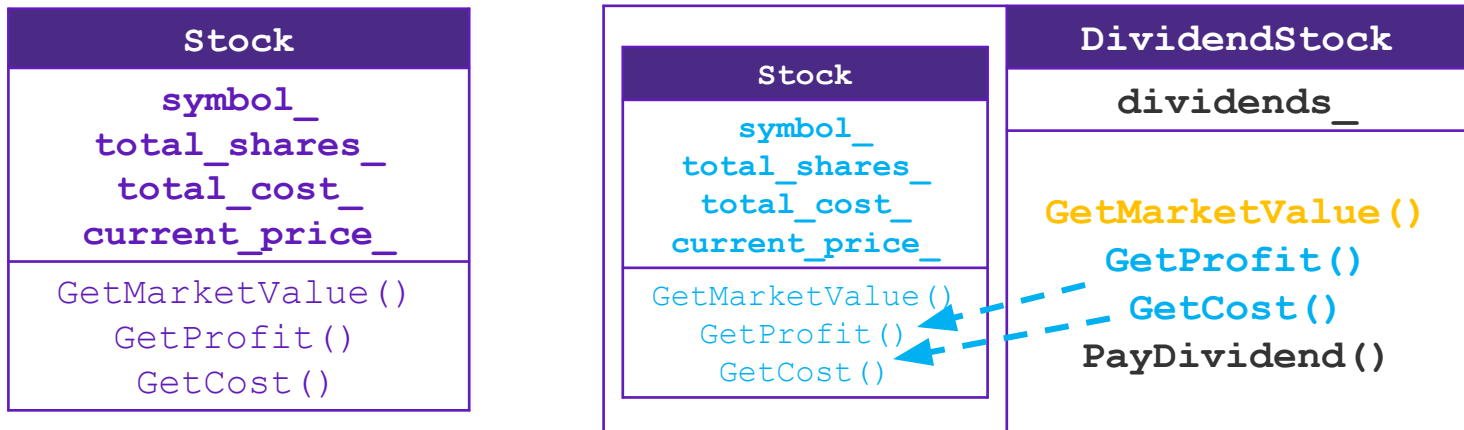
Stock
<code>symbol_</code> <code>total_shares_</code> <code>total_cost_</code> <code>current_price_</code>
<code>GetMarketValue()</code> <code>GetProfit()</code> <code>GetCost()</code>

BASE

DividendStock
<code>symbol_</code> <code>total_shares_</code> <code>total_cost_</code> <code>current_price_</code> <code>dividends_</code>
<code>GetMarketValue()</code> <code>GetProfit()</code> <code>GetCost()</code>

DERIVED

Back to Stocks



A derived class:

- **Inherits** the behavior and state (specification) of the base class
- **Overrides** some of the base class' member functions (optional)
- **Extends** the base class with new member functions, variables (optional)

Questions?



Polymorphism

& Dynamic Dispatch

Polymorphism in C++

```
PromisedType* var_p = new ActualType();
```

- `var_p` is a **pointer** to an object of `ActualType` on the Heap
- `ActualType` must be the same or a derived class of `PromisedType`
- `PromisedType` defines the *interface* (i.e. what can be called on `var_p`), but `ActualType` may determine which *version* gets invoked

Analogy: A box labeled “cell phone” could hold Android or iPhone

- `PromisedType` is the box, `ActualType` is the Android or iPhone

Dynamic Dispatch (like Java)

Usually, when a derived function is available for an object, we want the derived function to be invoked

- This requires a **run time** decision of what code to invoke
- This is the behavior in Java

A member function invoked on an object should be the **most-derived** *function* accessible to the object's visible type

- Can determine what to invoke from the **object** itself

Is this a Stock or a DividendStock ?

Example: `void PrintStock (Stock* s) { s->Print (); }`

- Calls the appropriate `Print ()` without knowing the exact class of `*s`, other than it is some sort of Stock

Requesting Dynamic Dispatch

Prefix the member function declaration with the `virtual` keyword

- Derived/child functions don't need to repeat `virtual`, but is traditionally good style to do so
- This is how method calls work in Java (no virtual keyword needed)
- You almost always want functions to be virtual
- C++ doesn't do dynamic dispatch by default so virtual keyword is strictly required if we want to make sure we're calling the most derived version of a function.

`override` keyword (C++11)

- Tells compiler this method should be overriding an inherited virtual function – **always** use when you can
- Prevents overloading vs. overriding bugs

Dynamic Dispatch Example

When a member function is invoked on an object:

- The *most-derived function* accessible to the object's visible type is invoked (decided at **run time** based on actual type of the object)

```
double DividendStock::GetMarketValue() const {  
    return get_shares() * get_share_price() + dividends;  
}
```

```
double DividendStock::GetProfit() const {    // inherited  
    return GetMarketValue() - GetCost();  
}      Should call DividendStock::GetMarketValue()
```

DividendStock.cc

```
double Stock::GetMarketValue() const {  
    return get_shares() * get_share_price();  
}
```

```
double Stock::GetProfit() const {  
    return GetMarketValue() - GetCost();  
}
```

Stock.cc

Inherited
from
stock



Dynamic Dispatch Example

```
#include "Stock.h"
#include "DividendStock.h"
```

```
DividendStock dividend();
DividendStock* ds = &dividend;
Stock* s = &dividend;    // why is this allowed?
```

A DividendStock “is-a” Stock, and has every part of Stock’s interface

```
// Invokes DividendStock::GetMarketValue()
ds->GetMarketValue();
```


```
// Invokes DividendStock::GetMarketValue()
s->GetMarketValue();
```

```
// invokes Stock::GetProfit(), since that method is inherited.
// Stock::GetProfit() invokes DividendStock::GetMarketValue(),
// since that is the most-derived accessible function.
s->GetProfit();
```

Most-Derived

```
class A {  
    public:  
        // Foo will use dynamic dispatch  
        virtual void Foo();  
};  
  
class B : public A {  
    public:  
        // B::Foo overrides A::Foo  
        virtual void Foo();  
};  
  
class C : public B {  
    // C inherits B::Foo()  
};
```


```
void Bar() {  
    A* a_ptr;  
    C c;  
  
    a_ptr = &c;  
  
    // Whose Foo() is called?  
    a_ptr->Foo();  
}
```



Most-Derived

```
class A {  
    public:  
        // Foo will use dynamic dispatch  
        virtual void Foo();  
};  
  
class B : public A {  
    public:  
        // B::Foo overrides A::Foo  
        virtual void Foo();  
};  
  
class C : public B {  
    // C inherits B::Foo()  
};
```

```
void Bar() {  
    A* a_ptr;  
    C c;  
  
    a_ptr = &c;  
  
    // B::Foo() is called  
    a_ptr->Foo();  
}
```





Poll Question (Pollev.com/cs374)

Whose **Foo()** is called?

- | | Q1 | Q2 |
|----|----|----|
| A. | A | B |
| B. | A | D |
| C. | B | B |
| D. | B | D |

```
void Bar() {  
    A* a_ptr;  
    C c;  
    E e;  
  
    // Q1:  
    a_ptr = &c;  
    a_ptr->Foo();  
  
    // Q2:  
    a_ptr = &e;  
    a_ptr->Foo();  
}
```

```
class A {  
    public:  
        virtual void Foo();  
};  
  
class B : public A {  
    public:  
        virtual void Foo();  
};  
  
class C : public B {  
};  
  
class D : public C {  
    public:  
        virtual void Foo();  
};  
  
class E : public C {  
};
```

W Who Foo() is called?

0

A B

0%

A D

0%


B B

0%


B D

0%





Poll Question (Pollev.com/cs374)



Has `Foo` definition

```
void Bar() {  
    A* a_ptr;  
    C c;  
    E e;  
  
    // Q1:  
    a_ptr = &c;  
    a_ptr->Foo();  
  
    // Q2:  
    a_ptr = &e;  
    a_ptr->Foo();  
}
```

```
class A {  
    public:  
        virtual void Foo();  
};  
  
class B : public A {  
    public:  
        virtual void Foo();  
};  
  
class C : public B {  
};  
  
class D : public C {  
    public:  
        virtual void Foo();  
};  
  
class E : public C {  
};
```

Questions?



How Can This Possibly Work?

The compiler produces `Stock.o` from *just* `Stock.cc`

- It doesn't know that `DividendStock` exists during this process
- So then how does the emitted code know to call

`Stock::GetMarketValue()` or `DividendStock::GetMarketValue()`

or something else that might not exist yet?

- **Function pointers!**

Stock.h

```
virtual double Stock::GetMarketValue() const;  
virtual double Stock::GetProfit() const;
```

```
double Stock::GetMarketValue() const {  
    return get_shares() * get_share_price();  
}
```

```
double Stock::GetProfit() const {  
    return GetMarketValue() - GetCost();  
}
```

Stock.cc

Virtual Table

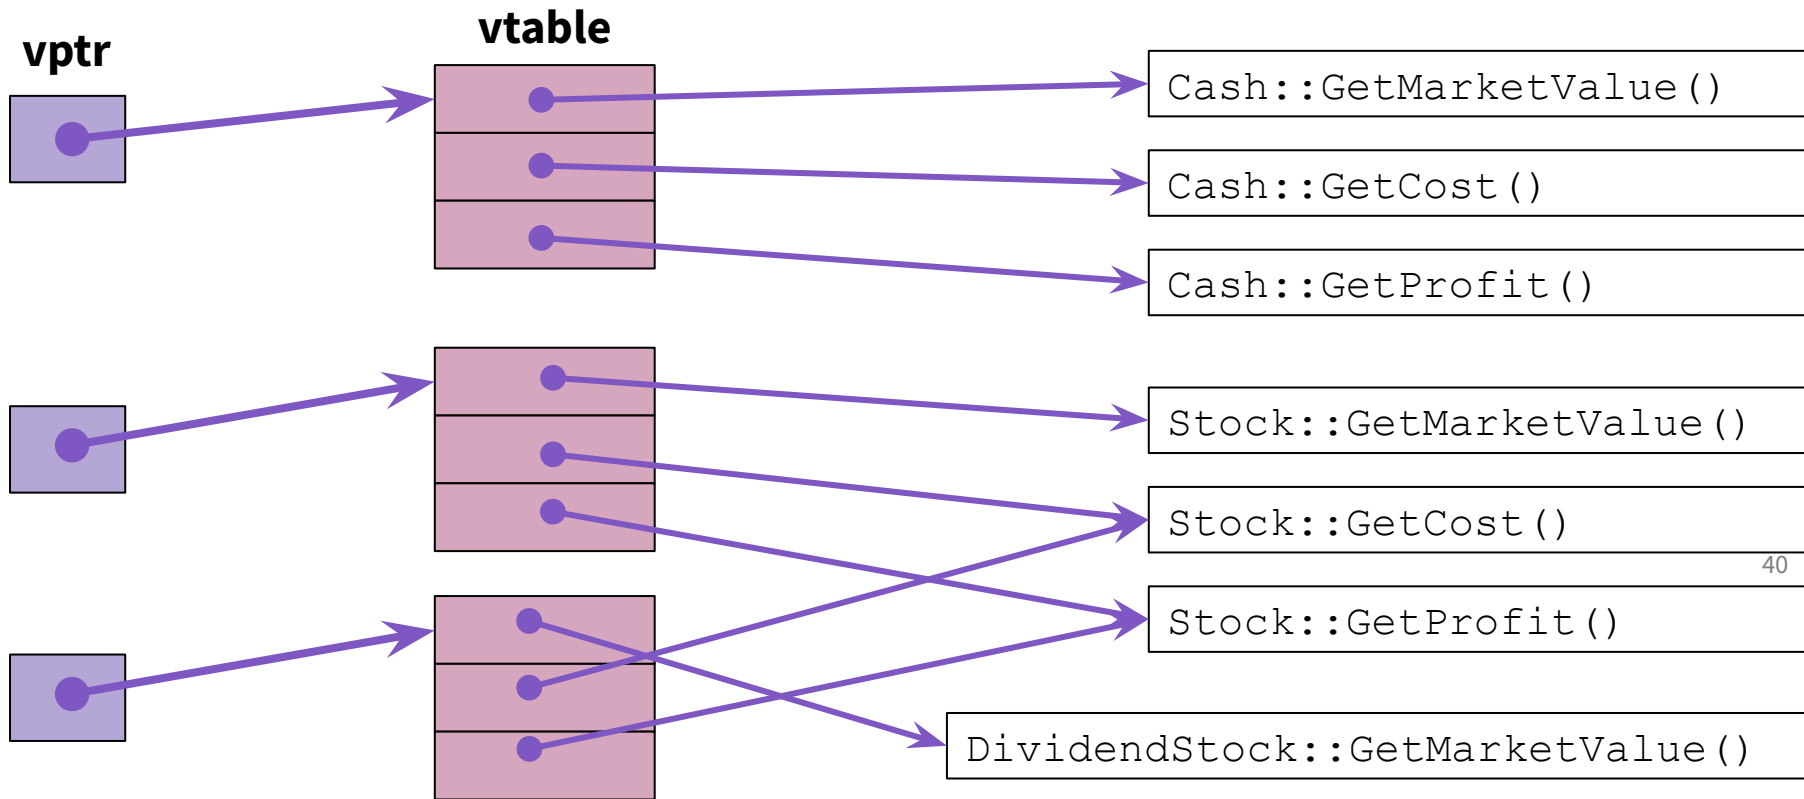
Virtual tables & virtual table
pointers

vtables and the vptr

If a class contains *any* virtual methods, the compiler emits:

- A (single) virtual function table (**vtable**) for *the class*
 - Contains a function pointer for each virtual method in the class
 - The pointers in the vtable point to the **most-derived** function for that class
- A virtual table pointer (**vptr**) for *each object instance*
 - A pointer to a virtual table as a “hidden” member variable
 - When the object’s constructor is invoked, the vptr is initialized to point to the vtable for the newly constructed object’s class
 - Thus, the vptr “remembers” what class the object is

vptr and vtable Visualization



vtable/vptr Example

```
class Base {  
    public:  
        virtual void f1();  
        virtual void f2();  
};  
  
class Der1 : public Base {  
    public:  
        virtual void f1();  
};  
  
class Der2 : public Base {  
    public:  
        virtual void f2();  
};
```

```
Base b;  
Der1 d1;  
Der2 d2;  
  
Base* b0ptr = &b;  
Base* b1ptr = &d1;  
Base* b2ptr = &d2;  
  
b0ptr->f1();  
b0ptr->f2();  
  
b1ptr->f1();  
b1ptr->f2();  
  
d2.f1();  
b2ptr->f1();  
b2ptr->f2();
```

vtable/vptr Example

```
class Base {  
    public:  
        virtual void f1();  
        virtual void f2();  
};  
  
class Der1 : public Base {  
    public:  
        virtual void f1();  
};  
  
class Der2 : public Base {  
    public:  
        virtual void f2();  
};
```

```
Base b;  
Der1 d1;  
Der2 d2;
```

```
Base* b0ptr = &b;  
Base* b1ptr = &d1;  
Base* b2ptr = &d2;
```

b0ptr->f1();	Base::f1()
b0ptr->f2();	Base::f2()

b1ptr->f1();	Der1::f1()
b1ptr->f2();	Base::f2()

d2.f1();	Base::f1()
b2ptr->f1();	Base::f1()
b2ptr->f2();	Der2::f2()

Static Dispatch

—

What happens if we omit “virtual”?

By default, without `virtual`, methods are dispatched **statically**

- At compile time, the compiler writes in a `call` to the address of the class' method based on the compile-time visible type of the callee
- This is *different* than Java

```
class Derived : public Base { ... };  
int main(int argc, char** argv) {  
    Derived d;  
    Derived* dp = &d;  
    Base* bp = &d;  
    dp->foo();  
    bp->foo();  
    return 0;  
}
```

Derived::foo()

...

Base::foo()

...

Static Dispatch Example

Removed `virtual` on methods:

Stock.h

```
double Stock::GetMarketValue() const;  
double Stock::GetProfit() const;
```

```
DividendStock dividend();  
DividendStock* ds = &dividend;  
Stock* s = &dividend;
```

```
ds->GetMarketValue(); // Calls DividendStock::GetMarketValue()  
s->GetMarketValue(); // Calls Stock::GetMarketValue()
```

```
s->GetProfit(); // Calls Stock::GetProfit(). Stock::GetProfit()  
calls Stock::GetMarketValue().
```

```
ds->GetProfit(); // Calls Stock::GetProfit(), since that method is  
inherited. Stock::GetProfit() calls Stock::GetMarketValue().
```

virtual is “sticky”

If `X::f()` is declared `virtual`, then a vtable will be created for class `X` and for **all** of its subclasses

- The vtables will include function pointers for (the correct) `f`

`f()` will be called using dynamic dispatch even if overridden in a derived class without the `virtual` keyword

- Good style to help the reader *and avoid bugs* by using `override`
 - Style guide controversy, if you use `override` should you use `virtual` in derived classes? Recent style guides say just use `override`, but you’ll sometimes see both, particularly in older code

Why Not Always Use `virtual`?

Two (fairly uncommon) reasons:

- Efficiency:
 - Non-virtual function calls are a tiny bit faster (no indirect lookup)
 - A class with zero virtual functions has objects without a `vptr` field
- Control:
 - If `f()` calls `g()` in class `X` and `g` is not virtual, we're guaranteed to call `X::g()` and not `g()` in some subclass
 - Particularly useful for framework design

In Java, all methods are virtual, except `static` class methods, which aren't associated with objects

In C++, you can pick what you want

- Omitting `virtual` can cause obscure bugs
- (Most of the time, you want member function to be virtual)

Abstract Classes

Sometimes we want to include a function in a class but *only* implement it in derived classes

- In Java, we would use an abstract method
- In C++, we use a “pure virtual” function

- Example: `virtual string noise() = 0;`

A class containing *any* pure virtual methods is **abstract**

- You can't create instances of an abstract class
- Extend abstract classes and override methods to use them

A class containing *only* pure virtual methods is the same as a Java interface

- Pure type specification without implementations (e.g. `asset`)

Questions?



Walkthrough: HW8



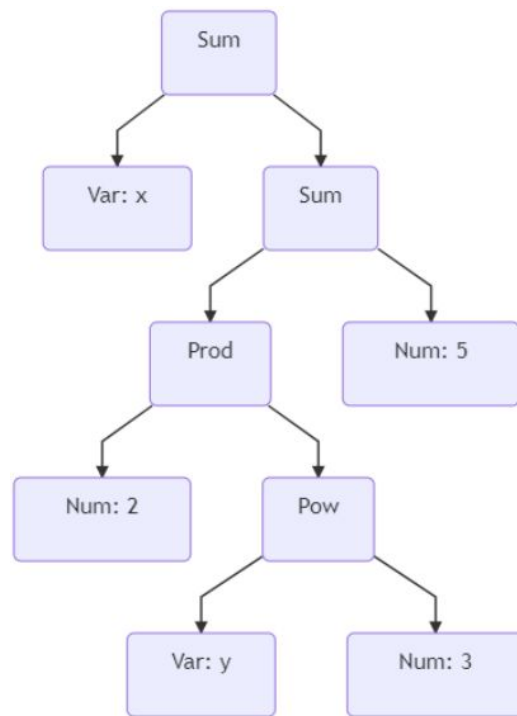
Abstract Syntax Tree

An AST is tree data structure that is commonly used to implement programming languages.

- The compiler parses your code line by line and organizes it into a tree of elements.
- In combination, these elements provide **semantic meaning**.


In HW8, you will not need to build this structure yourself.

- You will be implementing different types of `Expr` (via inheritance).




In HW8, each expression is mathematical.

$$2 * 7 + 3$$



$$2 * (7 + 3)$$



Evaluating Expressions

You will define classes for `Sum`, `Prod`, `Var`, and `Num`.

- These classes can be combined and used to ***evaluate*** expressions.

For example,

- `Prod(Sum(Num(5), Num(10)), Num(10)) == 150`
- `(5 + 10) * 10 == 150`

Generically, this same example could just be:

- `Expr(Expr(Expr, Expr), Expr)`

These expressions are naturally recursive (just like trees in general).

Functions like a calculator that accepts parameters (e.g. `X=5, Y=10`)

Ex21 due Friday, HW8 due Sunday!

Ex21 is due before the beginning of the next lecture

- Link available on the website:

<https://courses.cs.washington.edu/courses/cse374/24wi/exercises/>

HW8 due Sunday 11.59pm

- Instructions on course website:

<https://courses.cs.washington.edu/courses/cse374/24wi/homeworks/hw8/>

HW9 will be released Sunday and due the following Sunday

- You're almost to the end! 🎉