

CSE 374 Programming concepts and tools

Winter 2024

Instructor: Alex McKinney



Today

Assembly

- The software/hardware interface
- What actually goes on under the hood

Architecture Sits at the Hardware Interface

Source code

Different applications
or algorithms

Compiler


Perform optimizations,
generate instructions

Architecture

Instruction set

Hardware

Different
implementations



Demo: Compiler Explorer (godbolt.org)

Definitions

Instruction Set Architecture (ISA): The parts of a processor design that one needs to understand to write assembly code

- “What is directly visible to software”
- The interface used to connect software to hardware
- Like a header file (**.h**) in C!


Microarchitecture: Implementation of the ISA

- “Micro” because it deals with the internal, smaller-scale aspects.
- CSE/EE 469

Instruction Set Architectures

The ISA defines:

- The system's **state** (e.g. registers, memory)
- The **instructions** the CPU can execute
- The **effect** that each of these instructions will have on the system state



Put simply, how instructions are encoded, executed, and manipulated by the processor.

Instructions

Your CPU executes instructions using **machine code**

- Machine code is a binary encoding of instructions

The human readable version of machine code is called **assembly**

What `gcc` does

- C code is **compiled** into assembly
- Assembly is **assembled** into machine code (inside of object files)
 - By the assembler
- Machine code in object files are **linked** into an executable

Instruction Set Philosophies

Complex Instruction Set Computing (CISC): Add more and more elaborate and specialized instructions as needed

- Lots of tools for programmers to use, but hardware must be able to handle all instructions
- x86-64 is CISC, but only a small subset of CISC instructions encountered with Linux programs

Reduced Instruction Set Computing (RISC): Keep instruction set small and regular

- Easier to build fast hardware
- Let software do the complicated operations by composing simpler ones

RISC-V is an open-source ISA growing in popularity (used by NVIDIA)


RISC vs CISC

RISC puts more onerous on the compiler, but permits more optimizations.

- Finer granularity -> more room for reordering / parallelization.

The study of compiler optimizations is huge.

- Lots of graph algorithms applied in practice (e.g. [graph coloring](#)).
- A rich field of research.



CISC simplifies the compiler's job because the instruction handles so much more.

- Not as much room for optimization.
- Can't easily decompose the instruction into multiple steps.

RISC vs CISC

RISC

```
LOAD A, 2:3  
LOAD B, 5:2  
PROD A, B  
STORE 2:3, A
```

CISC

```
MULT 2:3, 5:2
```

Four instructions are consolidated into one!

Ref: [Stanford](#)

General ISA Design Decisions

Instructions

- What instructions are available? What do they do?
- How are they encoded?

Registers


- How many registers are there?
- How wide are they (e.g. 8-bit, 16-bit, etc)?
 - Often the same size as the address space (e.g. 32-bit, 64-bit).

Memory

- How do you specify a memory location?

Mainstream ISAs

Different architectures have different assembly syntax



x86

Designer	Intel, AMD
Bits	16-bit, 32-bit and 64-bit
Introduced	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
Design	CISC
Type	Register-memory
Encoding	Variable (1 to 15 bytes)
Endianness	Little

Old Macbooks & PCs
(Core i3, i5, i7, M)
x86-64 Instruction Set



ARM architectures

Designer	ARM Holdings
Bits	32-bit, 64-bit
Introduced	1985; 31 years ago
Design	RISC
Type	Register-Register
Encoding	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 user-space compatibility ^[1]
Endianness	Bi (little as default)

New Macbooks &
Smartphone-like devices
(iPhone, iPad, Raspberry Pi)
ARM Instruction Set



MIPS

Designer	MIPS Technologies, Inc.
Bits	64-bit (32→64)
Introduced	1981; 35 years ago
Design	RISC
Type	Register-Register
Encoding	Fixed
Endianness	Bi

Digital home & networking
equipment
(Blu-ray, PlayStation 2)
MIPS Instruction Set


Writing Assembly Code? In 2024?

Chances are, you'll never write a program in assembly, but understanding assembly is the key to the machine-level execution model:

- Behavior of programs in the presence of bugs
 - When high-level language model breaks down
- Tuning program performance
- Implementing systems software
- Fighting malicious software
 - Distributed software is in binary form

Like how C will help you better appreciate the magic that Java does for you. Similarly, understanding assembly will help you understand what the machine is doing for you.

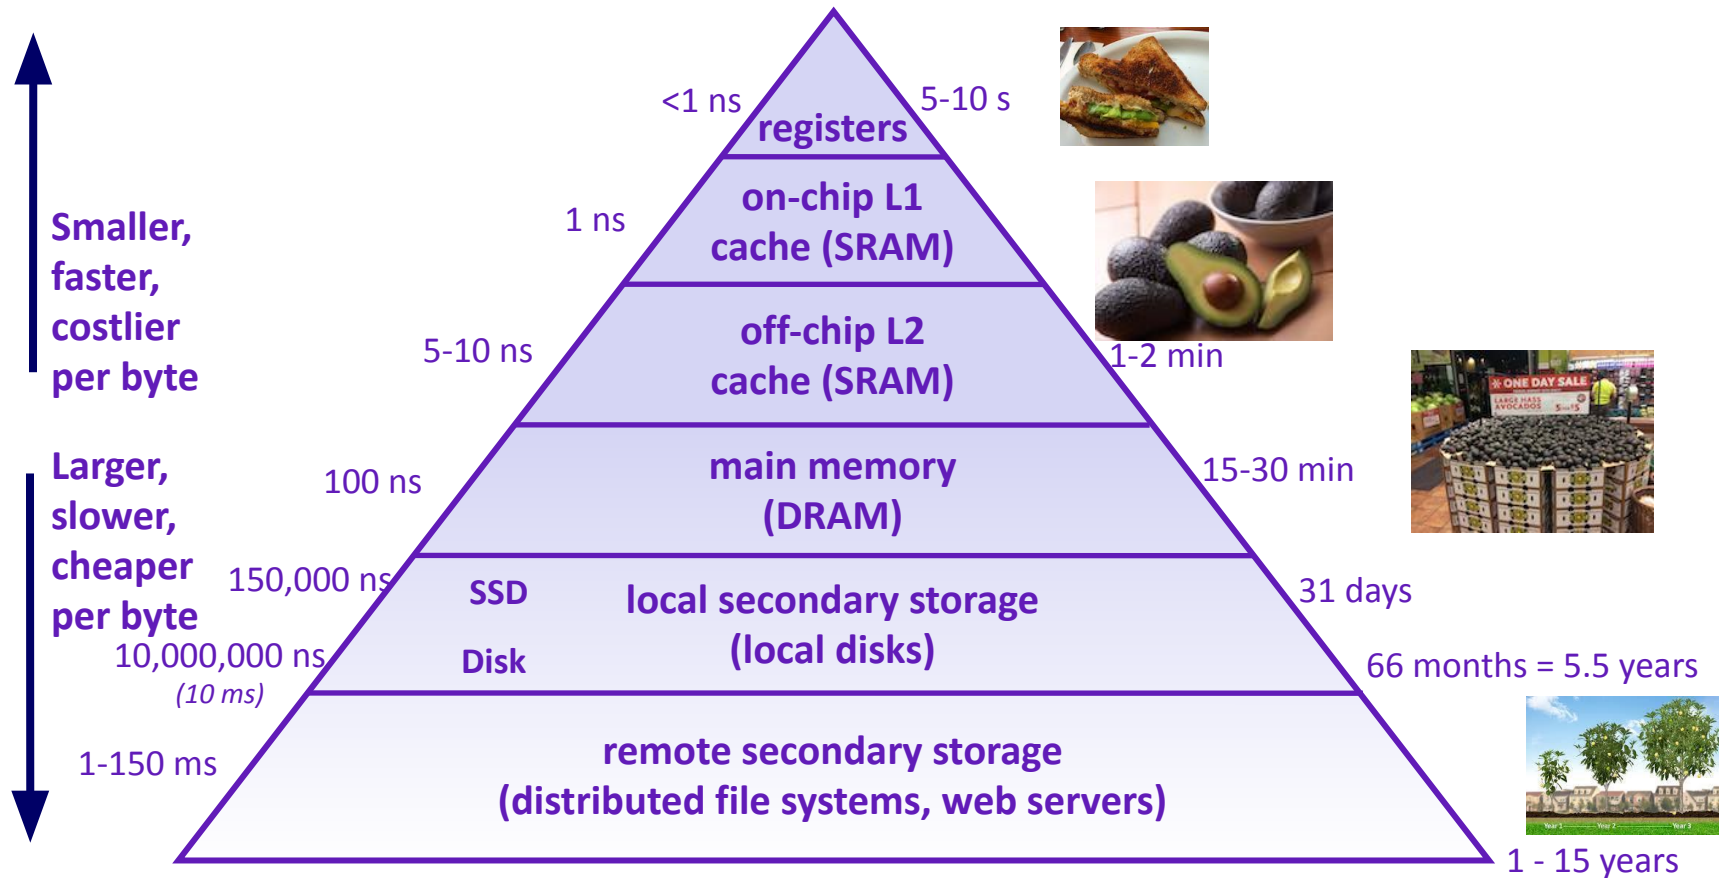
Assembly Programmer's View



Programmer-visible state

- **PC:** the Program Counter (`%rip` / Instruction Pointer in x86-64)
 - Address of next instruction
- **Named registers**
 - Heavily used program data
- **Condition codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

Review: Memory Hierarchy



What is a Register?

A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle)

Registers have *names*, not *addresses*

- In assembly, they start with `%` (e.g. `%rsi`)

Registers are at the heart of assembly programming

- They are a precious commodity in all architectures, but *especially* x86
- Complicated algorithms are used to ensure that registers are used wisely.

x86-64 Integer Registers – 64 bits wide

r = 8 bytes

e = 4 bytes

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

Stack
pointer

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

Memory

- Addresses
 - `0x7FFFD024C3DC`
- Big
 - ~ 8 GiB
- Slow
 - ~50-100 ns
- Dynamic
 - Can “grow” as needed while program runs

Registers

- Names
 - `%rdi`
- Small
 - $(16 \times 8 \text{ B}) = 128 \text{ B}$
- Fast
 - sub-nanosecond timescale
- Static
 - fixed number in hardware

Three Basic Kinds of Instructions

1. Transfer data between memory and register

- **Load** data from memory into register
 - `%reg = Mem[address]`
- **Store** register data into memory
 - `Mem[address] = %reg`

Remember: Memory is indexed just like an array of bytes!

2. Perform arithmetic operation on register or memory data

- `c = a + b; z = x << y; i = h & g;`

3. Control flow: what instruction to execute next

- Unconditional jumps to/from procedures
- Conditional branches

Operand types

Immediate: Constant integer data

- Examples: `$0x400`, `$-533`
- Like C literal, but prefixed with ``$'`
- Encoded with 1, 2, 4, or 8 bytes

Register: 1 of 16 integer registers

- Examples: `%rax`, `%r13`

Memory: Consecutive bytes of memory at a computed address

- Simplest example: `(%rax)`

`%rax`

`%rcx`

`%rdx`

`%rbx`

`%rsi`

`%rdi`

`%rsp`

`%rbp`

`%rN`

Moving Data

General form: `mov_ source, destination`

- Missing letter () specifies size of operands
- Lots of these in typical code

Example

- `movb src, dst`
 - Move 1-byte “byte”
- `movw src, dst`
 - Move 2-byte “word”
- `movl src, dst`
 - Move 4-byte “long word”
- `movq src, dst`
 - Move 8-byte “quad word”

Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4, %rax	rax = 4;
		Mem	movq \$-147, (%rax)	*rax = -147;
	Reg	Reg	movq %rax, %rdx	rdx = rax;
		Mem	movq %rax, (%rdx)	*rdx = rax;
	Mem	Reg	movq (%rax), %rdx	rdx = *rax;

() is like * in C

What about mem to mem?

Cannot do memory-memory transfer with a single instruction!

How would you do it?

1. Mem -> Reg
2. Reg -> Mem

Requiring separate steps is important:

- Actually increases opportunities for more optimizations.
- Keeps the ISA as simple as possible.
- More consistent and predictable.

Some Arithmetic Operations

Binary (two-operand) Instructions:

- Beware argument order!
- How do you implement “`r3 = r1 + r2`”?

Let say:

- `r1` is in `%rdi`
- `r2` is in `%rsi`
- `r3` is in `%rax`

```
movq %rdi, %rax
```

```
addq %rsi, %rax
```

Format	Computation	
<code>addq src, dst</code>	<code>dst = dst + src</code>	<code>(dst += src)</code>
<code>subq src, dst</code>	<code>dst = dst - src</code>	
<code>imulq src, dst</code>	<code>dst = dst * src</code>	signed mult
<code>shrq src, dst</code>	<code>dst = dst >> src</code>	
<code>shlq src, dst</code>	<code>dst = dst << src</code>	(same as <code>salq</code>)
<code>xorq src, dst</code>	<code>dst = dst ^ src</code>	
<code>andq src, dst</code>	<code>dst = dst & src</code>	
<code>orq src, dst</code>	<code>dst = dst src</code>	

q = operand size specifier
(e.g. b, w, l, q = 1, 2, 4, 8)

Some Arithmetic Operations

Unary (one-operand) Instructions:

Format	Computation	
incq <i>dst</i>	$dst = dst + 1$	increment
decq <i>dst</i>	$dst = dst - 1$	decrement
negq <i>dst</i>	$dst = -dst$	negate
notq <i>dst</i>	$dst = \sim dst$	bitwise complement

Arithmetic Example

```
long simple_arith(long x, long y)
{
    long t1 = x + y;
    long t2 = t1 * 3;
    return t2;
}
```

Register	Use(s)
<u>%rdi</u>	<u>1st argument (x)</u>
<u>%rsi</u>	<u>2nd argument (y)</u>
<u>%rax</u>	<u>return value</u>

```
y += x;
y *= 3;
long r = y;
return r;
```

```
simple_arith:
    addq    %rdi, %rsi
    imulq   $3,  %rsi
    movq    %rsi, %rax
    ret
```

Compiler knows

- Arithmetic ops are like +=
- And it doesn't need intermediate values anyway
- So it transforms it to...


Example of Basic Addressing Modes

```
void swap(long* xp, long* yp) {  
    long t0 = *xp;  
    long t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

```
swap:  
    movq    (%rdi), %rax  
    movq    (%rsi), %rdx  
    movq    %rdx, (%rdi)  
    movq    %rax, (%rsi)  
    ret
```

Understanding swap()

```
void swap(long* xp, long* yp) {  
    long t0 = *xp;  
    long t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```



```
swap:  
    movq    (%rdi), %rax  
    movq    (%rsi), %rdx  
    movq    %rdx, (%rdi)  
    movq    %rax, (%rsi)  
    ret
```

<u>Register</u>		<u>Variable</u>
%rdi	↔	xp
%rsi	↔	yp
%rax	↔	t0
%rdx	↔	t1

Understanding swap()

Registers

%rdi	0x120
%rsi	0x100
%rax	
%rdx	


Memory

Word Address
123
0x120
0x118
0x110
0x108
456
0x100

swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```


Understanding swap()



swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```


Understanding swap()



swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```


Understanding swap()



swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)  # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```


Understanding swap()



swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```

What about Java, Bash, etc.?

C is translated directly into assembly (and then into machine code)

Other languages may be translated into another form

- Java is translated into an assembly-like form, which is then run by the Java interpreter/runtime
- The Java runtime is executing assembly instructions!

Some languages are directly interpreted without being translated into another form

- Most Bash implementations will directly interpret the commands without compiling
- Python can do either. It can be used as an interpreter or compile scripts

The Hardware/Software Interface

What's after Assembly?

Every line of assembly gets **translated into binary**.

- The length of the instruction varies based on the system
- 16-bit, 32-bit, 64-bit.

Generally composed of the following:

- **Opcode** (add, multiply, move, etc).
- **Registers** (%r1, %r2)
- **Immediate values** (\$3)

Registers and immediate values are **operands**

```
00000000 0000 0001 0001 1010 0010 0001 0004 0128
00000010 0000 0016 0000 0028 0000 0010 0000 0020
00000020 0000 0001 0004 0000 0000 0000 0000 0000
00000030 0000 0000 0000 0010 0000 0000 0000 0204
00000040 0004 8384 0084 c7c8 00c8 4748 0048 e8e9
00000050 00e9 6a69 0069 a8a9 00a9 2828 0028 fdfe
00000060 00fc 1819 0019 9898 0098 d9d8 00d8 5857
00000070 0057 7b7a 007a bab9 00b9 3a3c 003c 8888
00000080 8888 8888 8888 8888 288e be88 8888 8888
00000090 3b83 5788 8888 8888 7667 778e 8828 8888
000000a0 d61f 7abd 8818 8888 467c 585f 8814 8188
000000b0 8b06 e8f7 88aa 8388 8b3b 88f3 88bd e988
000000c0 8a18 880c e841 c988 b328 6871 688e 958b
000000d0 a948 5862 5884 7e81 3788 1ab4 5a84 3eec
000000e0 3d86 dcb8 5cbb 8888 8888 8888 8888 8888
000000f0 8888 8888 8888 8888 8888 8888 8888 0000
0000100 0000 0000 0000 0000 0000 0000 0000 0000
*
0000130 0000 0000 0000 0000 0000 0000 0000 0000
000013e
```

Assembly <-> Binary

mov, \$r1, \$10

|

0001 001 0000000000001010

|

Op: Move | Register: r1 | Immediate: 10

Just like we've with other systems (e.g. networking), the computer knows how to interpret each set of bits.

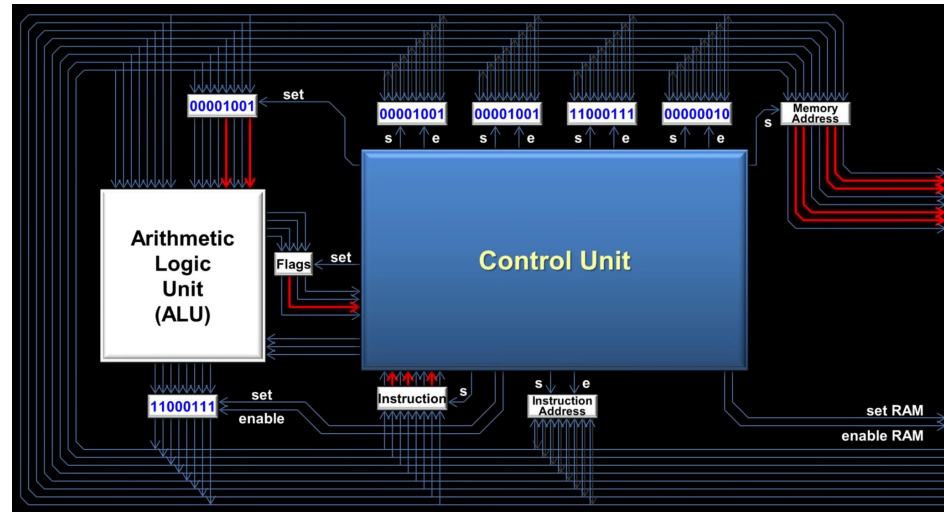
- It's a **protocol** (like floating point)!

CPU

The CPU parses the instruction string.

- Takes action based on what is encoded
- Interacts with other components as needed (e.g. ALU for math)

Fetch -> Decode -> Execute



The Big Picture


C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Terminal:



Assembly
language:

```
get_mpg:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    ...  
    popq     %rbp  
    ret
```

Shell
(Bash):




Machine
code:

```
0111010000011000  
1000110100000100000000010  
1000100111000010  
110000011111101000011111
```


OS:



Computer
system:



Computerphile



<https://www.youtube.com/watch?v=Z5JC9Ve1sfl&t=30s>